

# Reinforcement Learning

Alvin Li

December 2020

## 1 SARSA

- On-Policy
- Diff between Q-Learning: uses weighted average instead of max in Q-Learning
- $R + \gamma Q(S', A') - Q(S, A)$

## 2 Q-Learning

- Off-policy
- Learn policy to maximize total reward
- Q-table to store value at [state,action], updated using epsilon
- Algorithm:
  - (1) Agent at  $s_n$ , uses  $a_n$ , gets  $r_1$
  - (2) Pick action using either max value or random based on  $\epsilon$
  - (3) Update q
$$Q[\text{state}, \text{action}] = Q[\text{state}, \text{action}] + \text{lr} * (\text{reward} + \text{gamma} * \text{np.max}(Q[\text{new\_state}, :]) - Q[\text{state}, \text{action}])$$
- Markov Decision Process - outputs (rewards) are partly random (maybe due to malfunctioning) and partly under control, stochasticity. Also each state depends only on previous state.
- Bellman's Eqn

$$V(s) = \max_a (R(s, a) + \gamma V(s'))$$

- Bellman's Eqn with stochasticity

$$V(s) = \max_a (R(s, a) + \gamma \sum s' P(s, a, s') V(s'))$$

$P(s, a, s')$  is the probability of going from state  $s$  to  $s'$  given  $a$

- Q Eqn (quality of action)

$$Q(s, a) = R(s, a) + \gamma \sum_{s'} P(s, a, s') V(s')$$

$$Q(s, a) = R(s, a) + \gamma \sum_{s'} P(s, a, s') \max_{a'}(s', a')$$

- Temporal Difference

$$TD = R(s, a) + \gamma \sum_{s'} (P(s, a, s') \max_{a'}(s', a')) - Q(s, a)$$

- Update

$$Q(s, a) = Q_{t-1}(s, a) + \alpha TD$$

### 3 Double Deep Q-Learning


- Algorithm:
  - (1) Complete  $S, A, R, S_{t+1}$ ,  $a$  is chosen by  $\text{argmax}(Q)$  or epsilon greedy, store SARS in memory
  - (2) Sample a number of previous SARS
  - (3) Loss is MSE of pred  $Q$  vs target  $Q$ , using target  $Q$  ( $Q'$ ) for action selection, primary  $Q$  for evaluation  
ie.  $(r + \gamma Q(s_{t+1}, \text{argmax}_a Q'(s_{t+1}, a; \theta_t)) - Q(s, a, \theta))^2$
  - (4) Gradient Descent
  - (5) After  $N$  its, copy pred (or primary) network weights to target network weights
- Two NNs are used. One to estimate the target, one to estimate the pred (or primary). Target network is only updated every  $N$  its, using the prediction network.
- Experience Replay: sample a few  $[S, A, R, S]$  from previous  $X$  iterations to train
- Alternative: Update target network using Polyak Averaging,  $\tau$  is usually very small.  $\theta'$  are the weights of primary network,  $\theta$  is for the target network

$$\theta' \leftarrow \tau \theta + (1 - \tau) \theta'$$

- Double: To solve moving target problem, unstable nets
- Clipped: Use the minimum of  $Q, Q'$  for value evaluation, both using their own  $Q$  for action selection. To solve overestimation problem

$$r + \gamma \min(Q, Q')$$

- Dueling: Using Value and Advantage functions for Q. Aggregation is  $Q = V + A - \text{mean}(A)$



## 4 Deep Deterministic Policy Gradients

- $\theta^Q, \theta^\mu, \theta^{Q'}$  and  $\theta^{\mu'}$  represent Q network, deterministic policy func, target Q and target policy network weights respectively. Target nets, as before, are time-delayed versions of primary nets.
- Algorithm:

$Q(s,a)$  is critic,  $\mu(s|\theta^\mu)$  is actor,  $N$  is noise

```

for m episodes:
  for T timesteps:
    Select action based on  $\mu(s|\theta^\mu) + N$ 
    Store SARS
    Sample SARS batch for experience replay
    Update Critic using loss of:
       $L = \text{MSE}(r + \gamma Q'(\mu'|\theta^{Q'}) - Q(\theta^Q))$ 
    Update Actor using policy gradient:
       $\Delta_{\theta^\mu} J = \text{mean}(\Delta Q(\theta^Q) \Delta_{\theta^\mu} \mu(\theta^\mu))$ 
    Update target nets using soft updates ( $\tau$ )
  
```

- OU Process: Generate noise

## 5 Important Notes

Symbol or Vocab	Meaning
A or a	Action
R	Reward
$\alpha$	Step Size
$\pi_t(a)$	Policy (the probability of selecting a at t)
$Q_\pi$	State Action Value Function (expected return given state and action)

## 6 Sources Consulted

- [1] <https://towardsdatascience.com/simple-reinforcement-learning-q-learning-fcddc4b6fe56>
- [2] <https://arxiv.org/pdf/1509.06461.pdf>
- [3] <https://arxiv.org/abs/1509.02971>